

Principal Component Regression

Sriya Cheedella, Mia Shu

12/2/2019

Introduction

■ Motivation

When we have more than two covariates, multicollinearity impacts our model construction, parameter estimation, and prediction. In order to reduce its impact on our model, we reduce multicollinearity among variables by fitting the Principal Components.

■ Methodology

Break the collinear parts into uncorrelated smaller parts

Definitions

■ Multicollinearity

Multicollinearity exists among the predictor variables when these variables are correlated among themselves.

Example: weight and height; education level and salary

■ Confounding

The result of multicollinearity is often termed confounding: the situation when the correlation between two variables is aberrant due to a third variable included in the analysis.

Regression Model

Simple linear regression

$$Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i$$

Y_i Response at i th trial

β_0, β_1 Regression coefficients

X_i Predictors at i th trial

$\varepsilon_i \stackrel{iid}{\sim} \mathcal{N}(0, \sigma^2)$ Random error

Matrix Representation

Model $Y = X\beta + \varepsilon$

Residual $e_i = Y_i - \hat{Y}_i = Y - Xb$

b is the estimated vector of β

$$\begin{bmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_n \end{bmatrix} = \begin{bmatrix} 1 & X_1 \\ 1 & X_2 \\ \vdots & \vdots \\ 1 & X_n \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} + \begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \vdots \\ \varepsilon_n \end{bmatrix}$$

Multiple Linear Regression

Model

$$Y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \cdots + \beta_p X_{ip} + \varepsilon_i$$

$$\begin{bmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_n \end{bmatrix} = \begin{bmatrix} 1 & X_{11} & X_{12} & \cdots & X_{1p} \\ 1 & X_{21} & X_{22} & \cdots & X_{2p} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & X_{n1} & X_{n2} & \cdots & X_{np} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{bmatrix} + \begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \vdots \\ \varepsilon_n \end{bmatrix}$$

Diagnostic for Multicollinearity

- `ggpairs(X)& cor(X)`

Look for high pairwise correlation

- `vif(X)`

5-10 moderately high

< 10 extremely high

The Least Squares Estimator

y	$\frac{\partial y}{\partial x}$
\mathbf{Ax}	\mathbf{A}^T
$\mathbf{x}^T \mathbf{A}$	\mathbf{A}
$\mathbf{x}^T \mathbf{x}$	$2\mathbf{x}$
$\mathbf{x}^T \mathbf{Ax}$	$\mathbf{Ax} + \mathbf{A}^T \mathbf{x}$

$$RSS(b) = \sum e_i^2 = \mathbf{e}^T \mathbf{e}$$

$$= (\mathbf{Y} - \mathbf{Xb})^T (\mathbf{Y} - \mathbf{Xb})$$

$$= \mathbf{Y}^T \mathbf{Y} - \mathbf{Y}^T \mathbf{Xb} - \mathbf{b}^T \mathbf{X}^T \mathbf{Y} + \mathbf{b}^T \mathbf{X}^T \mathbf{Xb}$$

$$\frac{dRSS}{db} = -2\mathbf{X}^T \mathbf{Y} + 2\mathbf{X}^T \mathbf{Xb} = 0 \quad \mathbf{b} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

Variance Covariance Matrix

- Denoted $\sigma^2\{b\}$

- Estimate $\sigma^2 \rightarrow s^2 = MSE = \frac{\sum e_i^2}{n-2}$

$$\sigma^2\{b\} = \sigma^2 (X^\top X)^{-1}$$

- Multicollinearity

$(X^\top X)^{-1}$ close to singular

Sensitive to small perturbation \Rightarrow Unreliable parameter estimates

Geometrical Representation




Figure 1: img geom

$$M = I - X(X^T X)^{-1} X^T$$

$$e = MY = Y + \hat{Y} = Y + Xb\hat{Y} = HY = Xb$$

$$H = X(X^T X)^{-1} X^T \Rightarrow \hat{Y} \perp e$$

Spectral Decomposition

$$A = \lambda_1 u_1 u_1^\top + \lambda_2 u_2 u_2^\top + \cdots + \lambda_n u_n u_n^\top$$

where A is a square symmetric matrix

$$A = P D P^\top$$

P Orthonormal eigenvectors

D Diagonal matrix of eigenvalues

Principal Component Analysis

- Reduce a large set of correlated predictor variables to a smaller uncorrelated set.
- The principal component for a set of vectors are a set of linear combinations of the vectors chosen so that such set captures the most information in a smaller subset of vectors.

Procedure

- Standardize $\frac{X - \mu}{\sigma}$
- Find $X^T X = P D P^T = Z^T Z$

Singular Value Decomposition of $X \Rightarrow$ Truncated SVD

Maximize Rayleigh Coefficients

$$w_1 = \operatorname{argmax} \left\{ \frac{w^T X^T X w}{w^T w} \right\}$$

$$x_k = x - \sum_{s=1}^{k-1} x w_s w_s^T$$

$$w_k = \operatorname{argmax} \left\{ \frac{w^T x_k^T x_k w}{w^T w} \right\}$$

Procedure Continued

- Step Two

Fit Y on Z (OLS)

- Step Three

Choose components

- Step Four

Transform back to x scale

RidgeReg Data

```
X1 <- 1:18; X2 <- c(2,4,6,7,7,7,8,10,12,13,13,13,14,16,18,19,19,19); X3 <- c(1,2,4,3,2,1,1,1,2,3,4,5,6,7,8,9,10,11)
testdf <- data.frame(cbind(X1,X2,X3,Y))
testmod <- lm(Y~., data = testdf)
# vif(testmod)
testpcr <- pcr(Y~., data = testdf, scale=TRUE, validation = "CV")
cor(testdf)
```

```
##           X1           X2           X3           Y
## X1  1.00000000  0.9878415 -0.01505107  0.9855438
## X2  0.98784149  1.0000000  0.13381266  0.9955738
## X3 -0.01505107  0.1338127  1.00000000  0.1165387
## Y   0.98554384  0.9955738  0.11653870  1.0000000
```

```
summary(testpcr)
```

```
## Data:      X dimension: 18 3
## Y dimension: 18 1
## Fit method: svdpc
## Number of components considered: 3
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps
## CV           11.19    1.992    1.282    1.308
## adjCV         11.19    1.721    1.269    1.289
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps
## X      66.50    99.97   100.00
## Y      99.05    99.05    99.15
```

Iris Data

```
irismod <- lm(Sepal.Length~., data = iris)
# vif(irismod)
cor(iris[1:4])
```

```
##              Sepal.Length Sepal.Width Petal.Length Petal.Width
## Sepal.Length    1.0000000   -0.1175698    0.8717538    0.8179411
## Sepal.Width     -0.1175698    1.0000000   -0.4284401   -0.3661259
## Petal.Length     0.8717538   -0.4284401    1.0000000    0.9628654
## Petal.Width      0.8179411   -0.3661259    0.9628654    1.0000000

irispcr <- pcr(Sepal.Length~., data = iris, scale = TRUE, validation = "CV")
summary(irispcr)
```

```
## Data:      X dimension: 150 5
## Y dimension: 150 1
## Fit method: svdpc
## Number of components considered: 5
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps
## CV           0.8308  0.5122  0.5086  0.3955  0.3348  0.3191
## adjCV        0.8308  0.5118  0.5080  0.3948  0.3341  0.3181
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps
## X           56.20   88.62   99.07   99.73  100.00
## Sepal.Length 62.71   63.58   78.44   84.95   86.73
```


Formula One Racing Data: Description

- Provides data from Formula One World Championships from 1950-2017 about constructors, lap times, race drivers, etc.
- Given 13 .csv files to parse from.
- We wanted to see which variables best captured the time spent on a circuit.
- The columns used were circuit times, number of laps, number of pit stops, pit stop times, constructor points and driver points.

Verify Multicollinearity

```
cor(hungres[, -c(1,2)])
```

```
##               milliseconds      laps      pitStops MillisecondsPS      conPoints
## milliseconds      1.00000000  0.92071476 -0.08693557   -0.03347315  0.03369049
## laps              0.92071476  1.00000000 -0.06366250   -0.06574144  0.33044566
## pitStops          -0.08693557 -0.06366250  1.00000000    0.97835968 -0.05373143
## MillisecondsPS    -0.03347315 -0.06574144  0.97835968    1.00000000 -0.15541410
## conPoints         0.03369049  0.33044566 -0.05373143   -0.15541410  1.00000000
## drivPoints        0.11927110  0.37747886 -0.10851265   -0.20113019  0.88398060
##               drivPoints
## milliseconds      0.1192711
## laps              0.3774789
## pitStops          -0.1085127
## MillisecondsPS    -0.2011302
## conPoints         0.8839806
## drivPoints        1.0000000
```

```
hungmod <- lm(milliseconds ~ laps + pitStops + MillisecondsPS + conPoints + drivPoints,
              data = hungres)
```

```
# vif(hungmod)
```

Principal Component Analysis

```
hungmat <- as.matrix(hungres[,-c(1,2)])  
pca <- prcomp(hungmat, scale = TRUE, center = TRUE)  
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6
## Standard deviation	1.5802	1.3618	1.2192	0.33964	0.19139	0.09921
## Proportion of Variance	0.4162	0.3091	0.2477	0.01923	0.00611	0.00164
## Cumulative Proportion	0.4162	0.7253	0.9730	0.99225	0.99836	1.00000

Principal Component Regression

```
hungpcr <- pcr(milliseconds ~ laps + pitStops + MillisecondsPS +  
               conPoints + drivPoints, data = hungres, scale = TRUE,  
               validation = "CV")  
summary(hungpcr)
```

```
## Data:      X dimension: 24 5  
## Y dimension: 24 1  
## Fit method: svdpc  
## Number of components considered: 5  
##  
## VALIDATION: RMSEP  
## Cross-validated using 10 random segments.  
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  
## CV      205719    207896    209380    78638    76811    75405  
## adjCV    205719    207447    208295    76846    75056    73201  
##  
## TRAINING: % variance explained  
##      1 comps  2 comps  3 comps  4 comps  5 comps  
## X      46.579    81.88    97.42    99.69    100.00  
## milliseconds  7.567    10.79    92.87    93.22    94.96
```

```
coef(hungpcr, intercept = TRUE)
```

```
## , , 5 comps  
##  
##      milliseconds  
## (Intercept)    1283658.573  
## laps          199457.435  
## pitStops      -157998.285  
## MillisecondsPS 154526.064  
## conPoints     -48469.888  
## drivPoints     5510.244
```

```
hungpcrlog <- pcr(log(milliseconds) ~ laps + pitStops +  
                  log(MillisecondsPS) + conPoints + drivPoints,  
                  data = hungres, scale = TRUE, validation = "CV")  
summary(hungpcrlog)
```

```
## Data:      X dimension: 24 5  
## Y dimension: 24 1  
## Fit method: svdpc  
## Number of components considered: 5  
##  
## VALIDATION: RMSEP  
## Cross-validated using 10 random segments.  
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  
## CV      0.03605    0.03608    0.03642    0.01537    0.01492    0.01311  
## adjCV    0.03605    0.03598    0.03624    0.01508    0.01464    0.01283  
##  
## TRAINING: % variance explained  
##      1 comps  2 comps  3 comps  4 comps  5 comps  
## X      46.592    81.87    97.40    99.66    100.00  
## log(milliseconds)  8.334    10.83    92.45    92.80    95.05
```

```
coef(hungpcrlog, intercept = TRUE)
```


```
## , , 5 comps  
##  
##      log(milliseconds)  
## (Intercept)    13.386699700  
## laps          0.035582342  
## pitStops      -0.029313474  
## log(MillisecondsPS) 0.028648215  
## conPoints     -0.009769960  
## drivPoints     0.001987028
```

Importance of Standardization

```
noscale <- prcomp(hungres[, -c(1,2)], scale = FALSE)
summary(noscale)

## Importance of components:
##          PC1          PC2          PC3          PC4          PC5          PC6
## Standard deviation 2.014e+05 1.081e+04 97.98 20.3 0.7074 0.08798
## Proportion of Variance 9.971e-01 2.870e-03 0.00 0.0 0.0000 0.00000
## Cumulative Proportion 9.971e-01 1.000e+00 1.00 1.0 1.0000 1.00000


biplot(noscale, scale = 0)
```



```
scale <- prcomp(hungmat, scale = TRUE, center = TRUE)
summary(scale)



## Importance of components:
##          PC1          PC2          PC3          PC4          PC5          PC6
## Standard deviation  1.5802  1.3618  1.2192  0.33964  0.19139  0.09921
## Proportion of Variance 0.4162  0.3091  0.2477  0.01923  0.00611  0.00164
## Cumulative Proportion 0.4162  0.7253  0.9730  0.99225  0.99836  1.00000

biplot(scale, scale = 0)
```





Validation Plots

```
validationplot(hungpcr, main = "Root Mean Squared Error") validationplot(hungpcr, val.type="MSEP", main =
```




Validation Plots (Log)

```
validationplot(hungpcrlog, main = "Root Mean ; validationplot(hungpcrlog, val.type="MSEP", ma
```



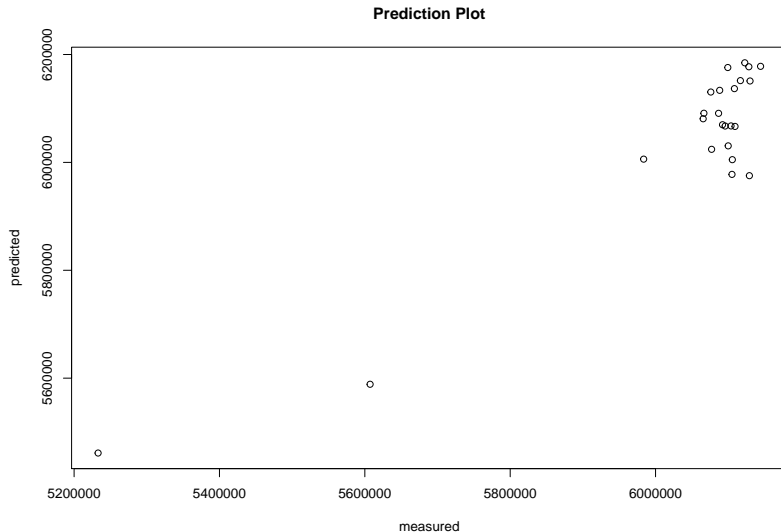
R^2 Plot

```
validationplot(hungpcr, val.type = "R2", main = "R2")
```



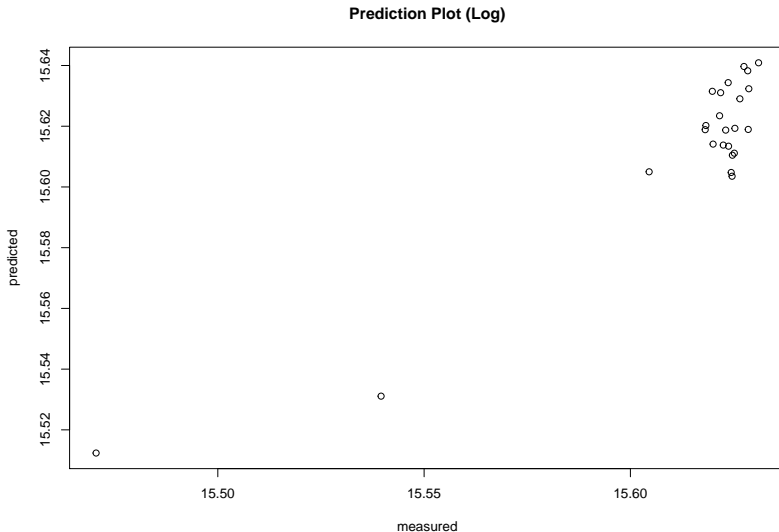
Prediction Plot

```
predplot(hungpcr, main = "Prediction Plot")
```



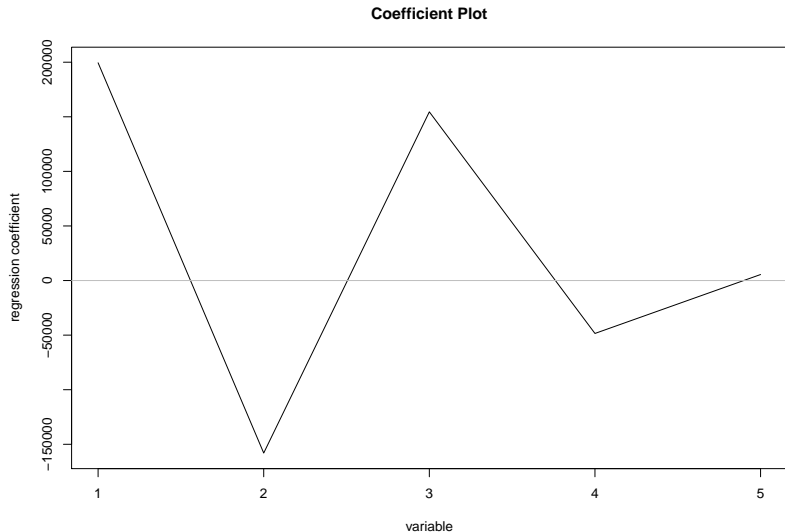
Prediction Plot (Log)

```
predplot(hungpcrlog, main = "Prediction Plot (Log)")
```



Coefficient Plot

```
coefplot(hungpcr, main = "Coefficient Plot")
```



References

- 1** Michael H. Kutner, Christopher J. Nachtsheim, and John Neter. Applied Linear Regression Models. New York, NY: McGraw-Hill/Irwin, 2004.
- 2** Johnston, R., Jones, K. & Manley, D. Qual Quant (2018) 52: 1957. <https://doi.org/10.1007/s11135-017-0584-6>
- 3** <https://www.whitman.edu/Documents/Academics/Mathematics/2017/Perez.pdf>
- 4** <https://datascienceplus.com/multicollinearity-in-r/>
- 5** https://web.njit.edu/~wguo/Math644_2012/Math644_Chapter%201_part2.pdf
- 6** https://ncss-wpengine.netdna-ssl.com/wp-content/themes/ncss/pdf/Procedures/NCSS/Principal_Components_Regression.pdf
- 7** https://en.wikipedia.org/wiki/Principal_component_analysis

References Continued

- 8** Dr. Christian Lucero, CMDA 4654 Lecture 08, Correlation and Least Squares, Lecture 14 MLR
- 9** Volodymyr Kuleshov, Fast algorithms for sparse principal component analysis based on Rayleigh quotient iteration
<http://proceedings.mlr.press/v28/kuleshov13.pdf>
- 10** https://ncss-wpengine.netdna-ssl.com/wp-content/themes/ncss/pdf/Procedures/NCSS/Principal_Components_Regression.pdf
- 11** <https://www.r-bloggers.com/performing-principal-components-regression-pcr-in-r/>